

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
```

In [3]:

```
bank = pd.read_csv('/Users/acer/Sandesh Pal/Data Science Assgn/LOgistic Regression/bank-full.csv', sep=';')
```

In [4]:

```
#Checking all the columns
bank.head()
```

Out[4]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	0	unknown
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	0	unknown
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1	0	unknown
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1	0	unknown
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1	0	unknown

In [5]:

```
bank.y.replace(('yes','no'), (1,0), inplace=True)
```

In [6]:

```
bank.default.replace(('yes','no'), (1,0), inplace=True)
```

In [7]:

```
bank.housing.replace(('yes','no'), (1,0), inplace=True)
```

In [8]:

```
bank.loan.replace(('yes','no'), (1,0), inplace=True)
```

In [9]:

```
bank.head(3)
```

Out[9]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome
0	58	management	married	tertiary	0	2143	1	0	unknown	5	may	261	1	-1	0	unknown
1	44	technician	single	secondary	0	29	1	0	unknown	5	may	151	1	-1	0	unknown
2	33	entrepreneur	married	secondary	0	2	1	1	unknown	5	may	76	1	-1	0	unknown

In [10]:

```
#Checking for na values
bank.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45211 non-null  int64
1   job         45211 non-null  object
2   marital     45211 non-null  object
3   education   45211 non-null  object
4   default     45211 non-null  int64
5   balance     45211 non-null  int64
6   housing     45211 non-null  int64
7   loan        45211 non-null  int64
8   contact     45211 non-null  object
9   day         45211 non-null  int64
10  month       45211 non-null  object
11  duration    45211 non-null  int64
12  campaign    45211 non-null  int64
13  pdays       45211 non-null  int64
14  previous    45211 non-null  int64
15  poutcome    45211 non-null  object
16  y           45211 non-null  int64
dtypes: int64(11), object(6)
memory usage: 5.9+ MB
```

In [11]:

```
bank = pd.get_dummies(bank)
```

In [12]:

```
bank.head(3)
```

Out[12]:

	age	default	balance	housing	loan	day	duration	campaign	pdays	previous	...	month_jun	month_mar	month_may	month_nov	month_c
0	58	0	2143	1	0	5	261	1	-1	0	...	0	0	1	0	
1	44	0	29	1	0	5	151	1	-1	0	...	0	0	1	0	
2	33	0	2	1	1	5	76	1	-1	0	...	0	0	1	0	

3 rows × 49 columns



In [13]:

```
bank.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 49 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                    45211 non-null  int64
1   default                               45211 non-null  int64
2   balance                               45211 non-null  int64
3   housing                               45211 non-null  int64
4   loan                                  45211 non-null  int64
5   day                                   45211 non-null  int64
6   duration                              45211 non-null  int64
7   campaign                              45211 non-null  int64
8   pdays                                45211 non-null  int64
9   previous                              45211 non-null  int64
10  y                                     45211 non-null  int64
11  job_admin.                            45211 non-null  uint8
12  job_blue-collar                       45211 non-null  uint8
13  job_entrepreneur                      45211 non-null  uint8
14  job_housemaid                         45211 non-null  uint8
15  job_management                        45211 non-null  uint8
16  job_retired                           45211 non-null  uint8
17  job_self-employed                     45211 non-null  uint8
18  job_services                          45211 non-null  uint8
19  job_student                           45211 non-null  uint8
20  job_technician                        45211 non-null  uint8
21  job_unemployed                        45211 non-null  uint8
22  job_unknown                           45211 non-null  uint8
23  marital_divorced                      45211 non-null  uint8
24  marital_married                       45211 non-null  uint8
25  marital_single                        45211 non-null  uint8
26  education_primary                     45211 non-null  uint8
27  education_secondary                   45211 non-null  uint8
28  education_tertiary                     45211 non-null  uint8
29  education_unknown                     45211 non-null  uint8
30  contact_cellular                      45211 non-null  uint8
31  contact_telephone                     45211 non-null  uint8
32  contact_unknown                       45211 non-null  uint8
33  month_apr                             45211 non-null  uint8
34  month_aug                             45211 non-null  uint8
35  month_dec                             45211 non-null  uint8
36  month_feb                             45211 non-null  uint8
37  month_jan                             45211 non-null  uint8
38  month_jul                             45211 non-null  uint8
39  month_jun                             45211 non-null  uint8
40  month_mar                             45211 non-null  uint8
41  month_may                             45211 non-null  uint8
42  month_nov                             45211 non-null  uint8
43  month_oct                             45211 non-null  uint8
44  month_sep                             45211 non-null  uint8
45  poutcome_failure                       45211 non-null  uint8
46  poutcome_other                         45211 non-null  uint8
47  poutcome_success                       45211 non-null  uint8
48  poutcome_unknown                       45211 non-null  uint8
dtypes: int64(11), uint8(38)
memory usage: 5.4 MB
```

In [14]:

```
bank[bank.duplicated()]
```

Out[14]:

```
age default balance housing loan day duration campaign pdays previous ... month_jun month_mar month_may month_nov month_oc
```

0 rows × 49 columns

```

In [15]:
```

```
#Dividing the dataset into X and Y variables
X = bank.loc[:,bank.columns!='y']
Y = np.ravel(bank.loc[:,bank.columns=='y'])
```

In [16]:

X

Out[16]:

	age	default	balance	housing	loan	day	duration	campaign	pdays	previous	...	month_jun	month_mar	month_may	month_nov	mor
0	58	0	2143	1	0	5	261	1	-1	0	...	0	0	1	0	
1	44	0	29	1	0	5	151	1	-1	0	...	0	0	1	0	
2	33	0	2	1	1	5	76	1	-1	0	...	0	0	1	0	
3	47	0	1506	1	0	5	92	1	-1	0	...	0	0	1	0	
4	33	0	1	0	0	5	198	1	-1	0	...	0	0	1	0	
...	
45206	51	0	825	0	0	17	977	3	-1	0	...	0	0	0	1	
45207	71	0	1729	0	0	17	456	2	-1	0	...	0	0	0	1	
45208	72	0	5715	0	0	17	1127	5	184	3	...	0	0	0	1	
45209	57	0	668	0	0	17	508	4	-1	0	...	0	0	0	1	
45210	37	0	2971	0	0	17	361	2	188	11	...	0	0	0	1	

45211 rows × 48 columns

```

In [17]:
```

Y

Out[17]:

```
array([0, 0, 0, ..., 1, 0, 0], dtype=int64)
```

In [18]:

```
#Building the logistic regression modellalit mehendiratta
model = LogisticRegression()
model.fit(X,Y)
```

```
C:\Users\acer\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Out[18]:

```
LogisticRegression()
```

In [19]:

```
#Predict for X dataset
y_pred = model.predict(X)
```

In [20]:

```
y_pred_df= pd.DataFrame({'actual': Y,
                          'predicted_prob': y_pred})
```

In [21]:

```
y_pred_df
```

Out[21]:

	actual	predicted_prob
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...
45206	1	1
45207	1	0
45208	1	1
45209	0	0
45210	0	0

45211 rows × 2 columns

In [22]:

```
# Confusion Matrix for the model accuracy
from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(Y,y_pred)
print (confusion_matrix)
```

```
[[39165  757]
 [ 4146 1143]]
```

In [23]:

```
#Classification report
from sklearn.metrics import classification_report
print(classification_report(Y,y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	39922
1	0.60	0.22	0.32	5289
accuracy			0.89	45211
macro avg	0.75	0.60	0.63	45211
weighted avg	0.87	0.89	0.87	45211

In [24]:


```
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

fpr, tpr, thresholds = roc_curve(Y, model.predict_proba (X)[:,1])

auc = roc_auc_score(Y, y_pred)

import matplotlib.pyplot as plt
plt.plot(fpr, tpr, color='red', label='logit model ( area = %0.2f)'%auc)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
plt.ylabel('True Positive Rate')
```

```
Text(0, 0.5, 'True Positive Rate')
```



```
auc
```

```
0.5985734647110931
```

Out[24]:



In [25]:

Out[25]:

In []: